

# Assignment 2

March 2021

## 1 Introduction

This assignment deals with understanding the scheduling of tasks in Linux with and without Preempt-RT. Preempt-RT is the official patch to make Linux more time predictable.

**Learning Objectives:** Understanding Linux scheduling.

### Expected deliverables.

You have to submit a report of 3 to 4 pages unless there are a lot of figures.

**Preparation:** In order to prepare for this lab, you have to install the following programs.

1. **KernelShark:** This is a kernel call stack tracing tool.
2. **Preempt-RT:** This is the official patch to make Linux more time predictable.

Following steps are tested on **Ubuntu 18.04 LTS (Long Term Support)**. The bold text represent the commands and the other text is to be followed in between steps. All of the referred figures are at the end of this document.

OR

Use the following virtual machine(VirtualBox)  
<https://surfdrive.surf.nl/files/index.php/s/2U46gdEh6uDNFHR>  
The username and password for login is: rts

### 1.1 Installing Preempt-RT Patch

Download the files **linux-5.4.19.tar.xz** and **patch-5.4.19-rt11.patch.xz** provided on the assignment page and place them in the same directory.

1. Extracting the Linux kernel.  
**tar xvf linux-5.4.19.tar.xz**
2. Enter the kernel directory.  
**cd linux-5.4.19**
3. Add patch to the kernel.  
**xzcat ../patch-5.4.19-rt11.patch.xz | patch -p1 -verbose**
4. Install Dependencies.  
**sudo apt-get install libncurses-dev libssl-dev**  
**sudo apt-get install bison**  
**sudo apt-get install flex**
5. Select Fully Preemptible Kernel (RT).  
**make menuconfig**
  - (i) Navigate using arrow keys to General Setup, then press Enter;
  - (ii) Look for Preemption Model, press Enter;
  - (iii) Choose Fully Preemptible Kernel (RT) option;
  - (iv) Return to the first screen by pressing Esc multiple times;
  - (v) Save .config file, then exit.
6. Building the kernel. (Take a cup of coffee and wait for it to accomplish)  
**make -j20**
7. Install the built modules.  
**sudo make modules\_install -j20**
8. Install the Patched Kernel  
**sudo make install -j20**  
If it asks to update Resume variable do steps 9 to 18 else jump to step 19.
9. Update the RESUME environment variable  
**echo "RESUME=UUID=8e8a8912-d6ee-4d7b-97bc-adab4a28eed9"**  
**| sudo tee /etc/initramfs-tools/conf.d/resume**  
In place of UUID replace with the UUID you see on your screen.  
**sudo update-initramfs -u -k all**
10. For Signing the signature of your patch, download the [mokconfig.cnf](#) file, and follow the steps below. To get a reference of the following steps click this link: <https://askubuntu.com/questions/1081472/vmlinuz-4-18-12-041812-generic-has-invalid-signature>.

11. Create the public and private key for signing the kernel  
**openssl req -config ./mokconfig.cnf -new -x509 -newkey rsa:2048 -nodes -days 36500 -outform DER -keyout "MOK.priv" -out "MOK.der"**
12. Convert the key also to PEM format (mokutil needs DER, sbsign needs PEM)  
**openssl x509 -in MOK.der -inform DER -outform PEM -out MOK.pem**
13. Enroll the key to your shim installation  
**sudo mokutil --import MOK.der**  
You will be asked for a password, you will just use it to confirm your key selection in the next step.
14. Restart your system  
Restart your system. You will encounter a blue screen of a tool called MOKManager. Select "Enroll MOK" and then "View key". Make sure it is your key you created. Afterwards continue the process and you must enter the password which you provided in previous step. Continue with booting your system.
15. Verify your key is enrolled via  
**sudo mokutil --list-enrolled**
16. Sign your installed kernel (it should be at /boot/vmlinuz-5.4.19-rt11)  
**sudo sbsign --key MOK.priv --cert MOK.pem /boot/vmlinuz-5.4.19-rt11 --output /boot/vmlinuz-5.4.19-rt11.signed**
17. Copy the initram of the unsigned kernel, so we also have an initram for the signed one.  
**sudo cp /boot/initrd.img-5.4.19-rt11{,signed}**
18. Update your grub-config  
**sudo update-grub**
19. Reboot your system and select the signed patched kernel (**5.4.19**) by going in to the Advanced boot options of Ubuntu shown in 1
20. When the system is up and running, enter the following command in your terminal:

**uname -a**

It should show you a message containing PREEMPT\_RT as shown in Figure 2.

Perform the following command to reduce the Round Robin scheduling timeslice.

**sudo sysctl kernel.sched\_rr\_timeslice\_ms=xx** where **xx**=milliseconds.

For this assignment, it is best to keep the value of **xx**=4 millisecond.

However, if you want to change the values in the experiment and want to explore more on your own, then feel free to increase its value and explain your observations.

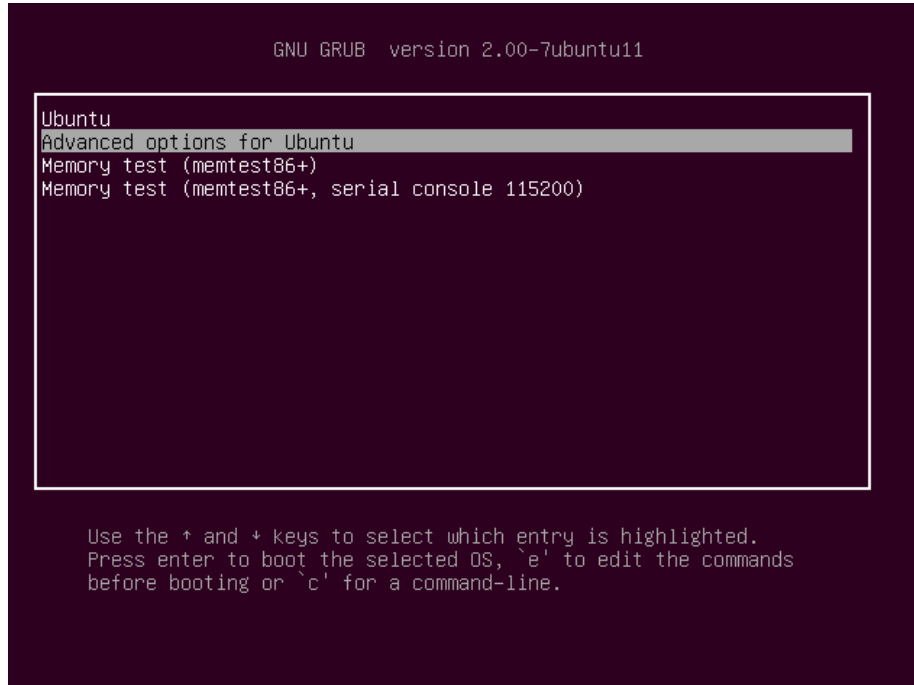
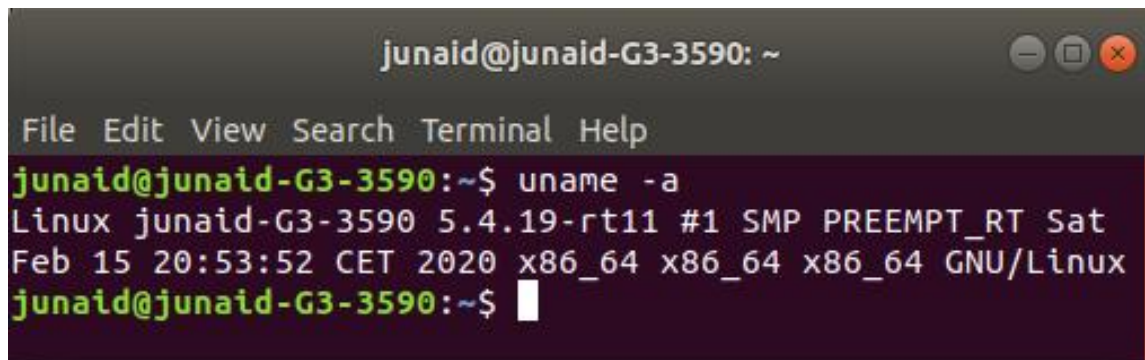


Figure 1: Advanced Options for Ubuntu

A terminal window titled 'junaid@junaid-G3-3590: ~' with standard window controls. The menu bar includes 'File', 'Edit', 'View', 'Search', 'Terminal', and 'Help'. The terminal shows the command 'junaid@junaid-G3-3590:~\$ uname -a' and its output: 'Linux junaid-G3-3590 5.4.19-rt11 #1 SMP PREEMPT\_RT Sat Feb 15 20:53:52 CET 2020 x86\_64 x86\_64 x86\_64 GNU/Linux'. The prompt 'junaid@junaid-G3-3590:~\$' is followed by a cursor.

```
junaid@junaid-G3-3590: ~
File Edit View Search Terminal Help
junaid@junaid-G3-3590:~$ uname -a
Linux junaid-G3-3590 5.4.19-rt11 #1 SMP PREEMPT_RT Sat
Feb 15 20:53:52 CET 2020 x86_64 x86_64 x86_64 GNU/Linux
junaid@junaid-G3-3590:~$
```

Figure 2: Preempt-RT Patch Confirmation

## 1.2 Installing KernelShark

Download the file [trace-cmd-kernelshark-v1.1.zip](#) provided on the assignment page.

1. Extract the contents of trace-cmd-kernelshark-v1.1.zip file.  
**cd trace-cmd-kernelshark-v1.1**
2. Build the source.  
**sudo make**
3. Build the graphical interface sources.  
**sudo make gui**
4. Install the built sources.  
**sudo make install**
5. Install the built graphical interface sources.  
**sudo make install\_gui**
6. If above steps do not make KernelShark appear in showapplications tab, then follow the instruction given on <https://kernelshark.org/build.html>

## 1.3 How to create a periodic thread?

In order to create periodic real-time tasks using POSIX, follow the tutorial in this link: <http://retis.sssup.it/~lipari/courses/str09/10.rtprogramming-handout.pdf>. For a general tutorial on how to create a POSIX based application,

follow the tutorial in this link: <https://computing.llnl.gov/tutorials/pthreads/>.

## 1.4 Creation and use of traces in KernelShark

This section explains the method to create a trace entry to be shown in KernelShark during the execution. Since KernelShark is only designed to show kernel call stacks, so it is important that any user level function or thread, which is required to be traced, makes the appropriate system calls. In the [FtraceExample.c](#) example source file provided with this assignment, in the [Assignment\\_2\\_Code\\_Template](#) folder, to get samples of thread creation, trace generation and CPU affinity setting.

The Figure 3 shows you the main of the example. Line 64 shows declaration of thread id variables for two threads. Line 66 opens a file from kernel debug facilities to insert markers in between program executions, which is another method of showing traces in KernelShark. If this method is not used then user has to invoke system calls, which does not mean anything except generation of a trace entry, which becomes cumbersome to manage. With this method user can insert his desired text in the traces along with variables values etc. It should be noticed that the file descriptor variable *trace\_fd* should be declared global so that all the threads can access it. This is followed by creation of threads, which execute the respective functions, and joining of threads. The function calls *trace\_write("JATW\_1")* and *trace\_write("JATW\_2")* are used to insert marker texts which appear in KernelShark as shown in Figure 4.

Figure 5 shows the code for one of the two threads, in which line number 16 to 25 shows the method to set CPU affinity. In this case, the thread has the affinity to CPU 1 which basically means that *Thread\_1* will only run on *CPU1*. The next step is the access to the *s\_mutex* for writing a trace message in to the *trace\_marker* file followed by rest of the operations. Figure 6 shows the simple function of formatting and writing a trace text in to the *trace\_marker*. **You can use provided code template for your assignment.**

**Command to create traces:** In order to show the trace of any program, it is important to first create the trace. Lets build the given [FtraceExample.c](#) example to create the trace first. The command to build this simple example is:

```
gcc -pthread FtraceExample.c -o aout
```

Once the program is built successfully, use the following command to generate the trace file.

```
sudo trace-cmd record -p function -F ./aout
```

With the above command, the trace will be generated by running the program. The **-F ./aout** option runs the program for which the traces will be generated. By default the name of the trace file will be **trace.dat**. If it is not required to observe the system and shared library functions calls in the trace, rather only your very own program traces then use the following command to generate the traces.

```
sudo trace-cmd record -p nop -F ./aout
```

In order to run the above command with an input argument to the built binary simply follow this command:

```
sudo trace-cmd record -p nop -F ./aout FIFO
```

The above command runs the built binary **./aout** with **SCHED\_FIFO** scheduling policy. Similarly, you can run it for **SCHED\_RR** and **SCHED\_OTHER** scheduling policies by passing **RR** and **OTHER** as arguments respectively.

**Showing traces in KernelShark:** Figure 7 shows an empty view of the KernelShark when it is opened. In this view, go to the top menu and select **File** → **Open** and then select the generated trace file which will show the view of Figure 8. From the drop down list of **Column** tab, select the option **Info** as it is shown in Figure 9. Then enter the text in the trace message search box next to the tab **contains** as it is shown in Figure 10 and press next to reach to the trace entry which is highlighted in green color. The timestamps column will provide the time at which that entry was written which will give an indication of the time when you recorded an event or start of a thread etc. Figure 9 also shows the orange and red colored bars against **CPU1** which shows the execution of **Thread\_1** and **Thread\_2** since their affinity was set to **CPU1** as it was shown in Figure 5.

Now that everything is set in the system, let's get on to the assignment itself. You have to use the **Assignment\_2\_Code\_Template.c** file which is provided in the **.zip** package.

## 2 Main questions

Answer the following questions in a digital format and upload it into the BrightSpace. Please do not forget to add your names, student numbers, and the date of submission to the file you upload.

**Question 1 (15 points).** Explain in detail how `SCHED_FIFO`, `SCHED_RR` and `SCHED_DEADLINE` work in Linux. You can learn about these policies by reading, for example, the following link: <http://man7.org/linux/man-pages/man7/sched.7.html>. Feel free to explore more on your own.

**Question 2 (30 points).** In order to understand the differences between the `SCHED_FIFO`, `SCHED_RR` and `SCHED_OTHER` scheduling policies, build a task set (by playing with `NUM_THREADS` and periods), schedule the task set with each of these policies, report your observations, compare the resulting schedules with each other, and explain the differences.

**Note 1.** Your answer must include explanation of what you have done, properties of the tasks you have generated (their WCET, period, priorities, utilization, etc.), the schedule visualized via KernelShark, and the KernelShark traces.

**Note 2.** To create a periodic task-set you can have a look at section 1.3. You need to build and run the program according to 1.4.

**Note 3.** Since we are using only a single CPU in this assignment, you must ensure that the total utilization of the task-set is not more than 1. You can use tracing functions to record traces whenever a thread starts or just before going to sleep to get a nice view of activities in KernelShark. Remember that it may incur some overhead, so it is up to you to decide on how you want to measure the response time. However, you must specify your decision in your report. Your report must also explain (briefly) how did you instrument the code for tracing.

**Note 4.** The second part of the slide 7, in this link <http://retis.sssup.it/~lipari/courses/str09/10.rtprogramming-handout.pdf>, provides a method to check the deadline miss of a task. The structure `next` is used to set the release time of the task and the structure `now` in the slide provides the finish time of the task. Similar to the deadline miss mechanism, you can take a difference between the finish time and release time of the task to measure the response time. Keep in mind that, to measure the response time you will have to use the previously set value of the `next` and not the new value which is set by the call `timespec_add_us`. You can add all of your measurements in an array and then dump it in a file, for the sake of analysis, at the end of your simulation.

**(30 points)** The following document explains three ways of building time-triggered tasks (see slide 5 till slide 7 of the <http://retis.sssup.it/~lipari/courses/str09/10.rtprogramming-handout.pdf>). Assume that our goal is to build periodic tasks. In that case, compare these three approaches by building a



task set (and scheduling it with rate-monotonic priorities using *SCHED\_FIFO*). On that task set, choose a task and then measure the difference between the expected arrival times of the task and the actual release times. Show that unlike the first two methods, the last one results in periodic activations. Justify your answer by adding the traces of KernelShark and visualizing the arrival times of the task

**Deadline:** The deadline of the assignment is **28 March, 2021 23:59**

```

62 int main(void)
63 {
64     pthread_t th1, th2;
65     printf("%s\n", "This is a Hellow World program to test trace_marker");
66     trace_fd = open("/sys/kernel/debug/tracing/trace_marker", O_WRONLY);
67     if(trace_fd < 0) {
68         printf("Error Opening trace marker, try running with sudo\n");
69     } else{
70         printf("Successfully opened trace_marker\n");
71     }
72
73     trace_write("JATW_1");
74
75     if(pthread_create(&th1, NULL, Thread1, NULL) != 0)
76     {
77         printf("Error Spawning Thread1");
78     }
79     else
80     {
81         printf("Spawned Thread1");
82     }
83
84     if(pthread_create(&th2, NULL, Thread2, NULL) != 0)
85     {
86         printf("Error Spawning Thread2");
87     }
88     else
89     {
90         printf("Spawned Thread2");
91     }
92
93     pthread_join(th1, NULL);
94     pthread_join(th2, NULL);
95
96     trace_write("JATW_2");
97
98     if(close(trace_fd) == 0)
99     {
100         printf("Successfully closed the file\n");
101     }
102
103     return 0;
104 }
105

```

Figure 3: main of the FtraceExample.c

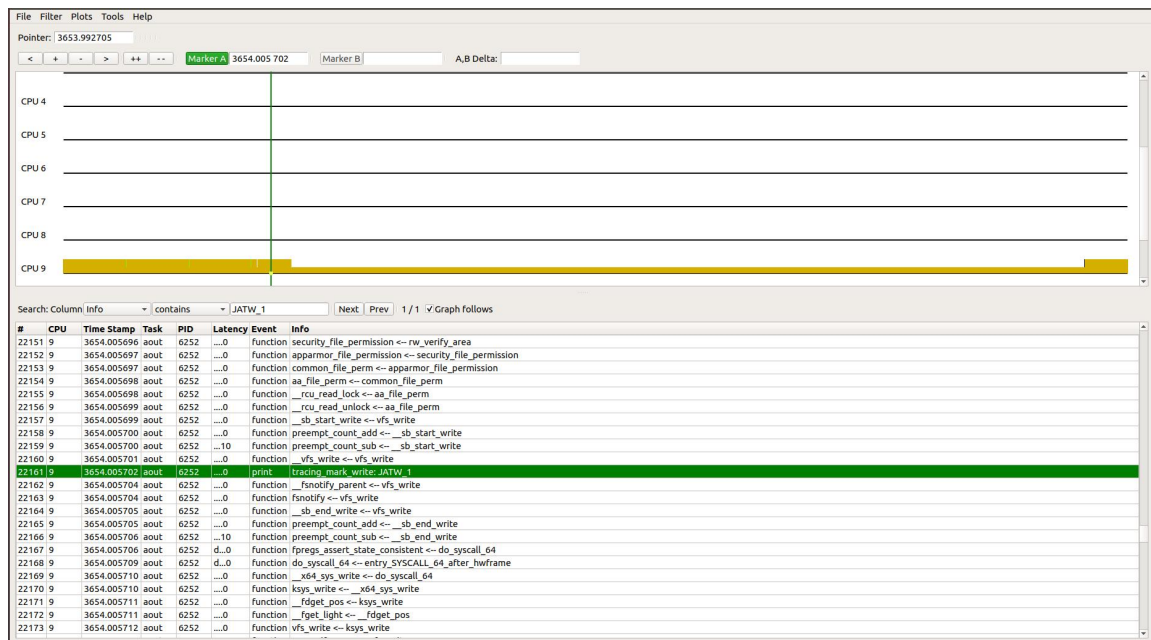


Figure 4: Trace marker JATW\_1

```

1  #define _GNU_SOURCE
2
3  #include <sys/fcntl.h>
4  #include <stdio.h>
5  #include <stdarg.h>
6  #include <pthread.h>
7  #include <sched.h>
8  #include <stdlib.h>
9  #include <unistd.h>
10
11 void trace_write(const char *fmt, ...);
12 static int trace_fd = -1;
13
14 void* Thread1(void *args)
15 {
16     cpu_set_t cpuset;
17     // CPU_ZERO: This macro initializes the CPU set set to be the empty set.
18     CPU_ZERO(&cpuset);
19     // CPU_SET: This macro adds cpu to the CPU set set.
20     CPU_SET(1, &cpuset);
21
22     const int set_result = pthread_setaffinity_np(pthread_self(), sizeof(cpu_set_t), &cpuset);
23     if (set_result != 0) {
24         printf("Error Setting CPU Affinity\n");
25     }
26
27     trace_write("JATW_Thread_1_1 ... Pi = %f", 3.14);
28
29     double testing = 0.0;
30     for(int a=0; a<1000000; a++)
31     {
32         testing = (a*testing)/(1+a);
33     }
34
35     trace_write("JATW_Thread_1_2 ... e = %f", 2.71);
36 }
37

```

Figure 5: Sample function Thread\_1

```

118 void trace_write(const char *fmt, ...)
119 {
120     va_list ap;
121     char buf[256];
122     int n;
123
124     if (trace_fd < 0)
125         return;
126
127     va_start(ap, fmt);
128     n = vsnprintf(buf, 256, fmt, ap);
129     va_end(ap);
130
131     write(trace_fd, buf, n);
132     printf("Written Successfully\n");
133 }
134

```

Figure 6: Sample function Trace\_Write

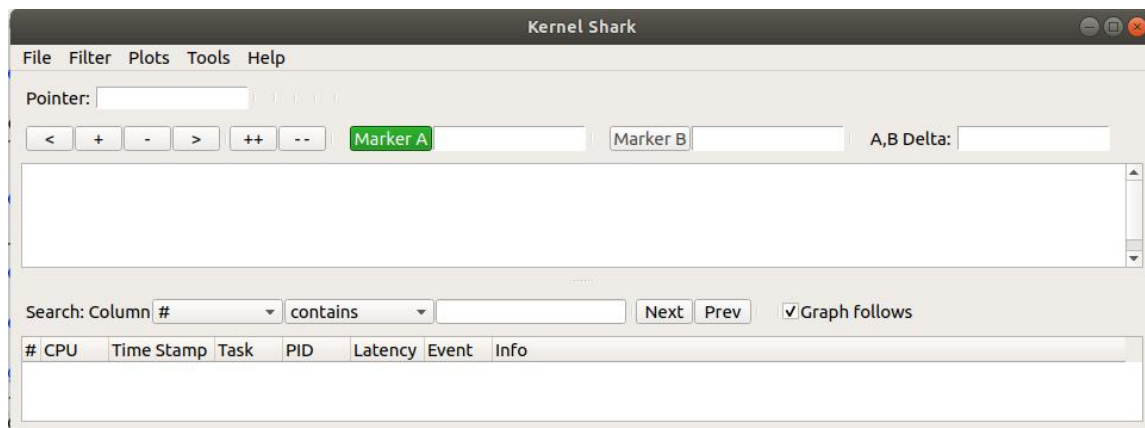


Figure 7: The empty view of KernelShark without any traces

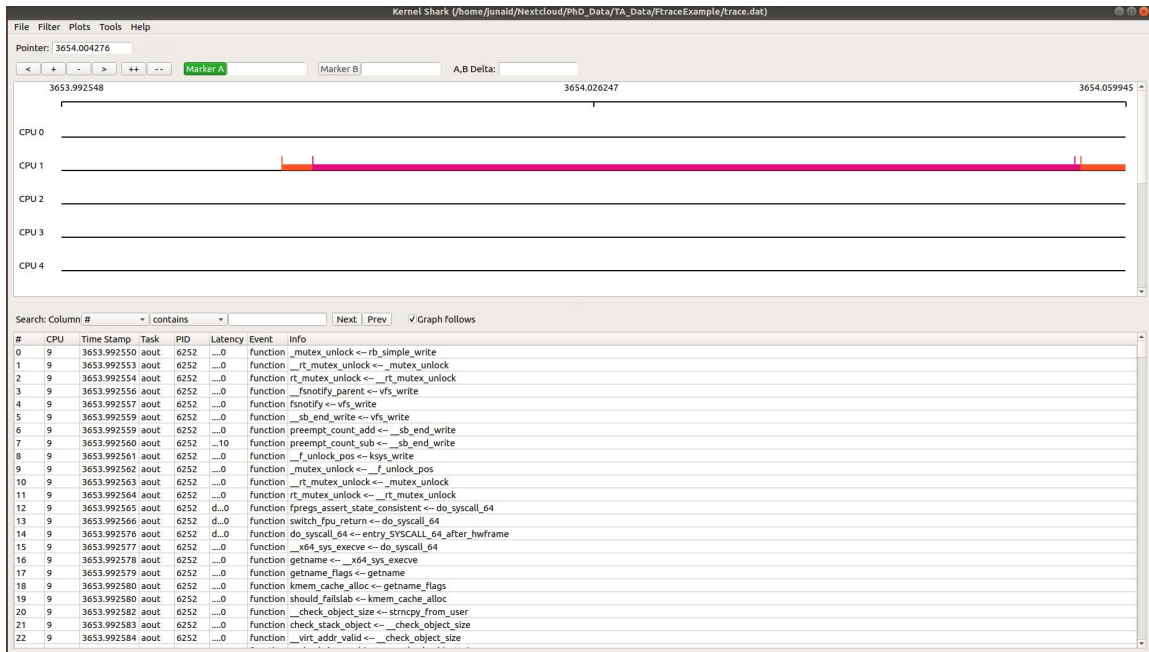


Figure 8: The view of KernelShark with traces

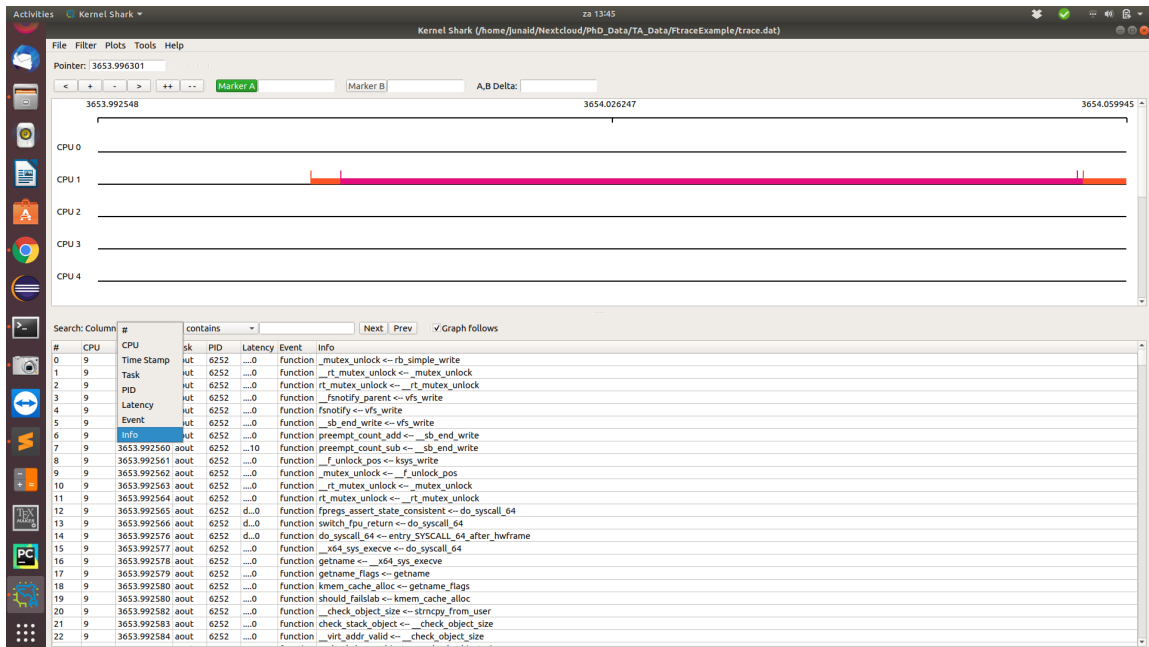


Figure 9: Select "Info" from the drop down list of column

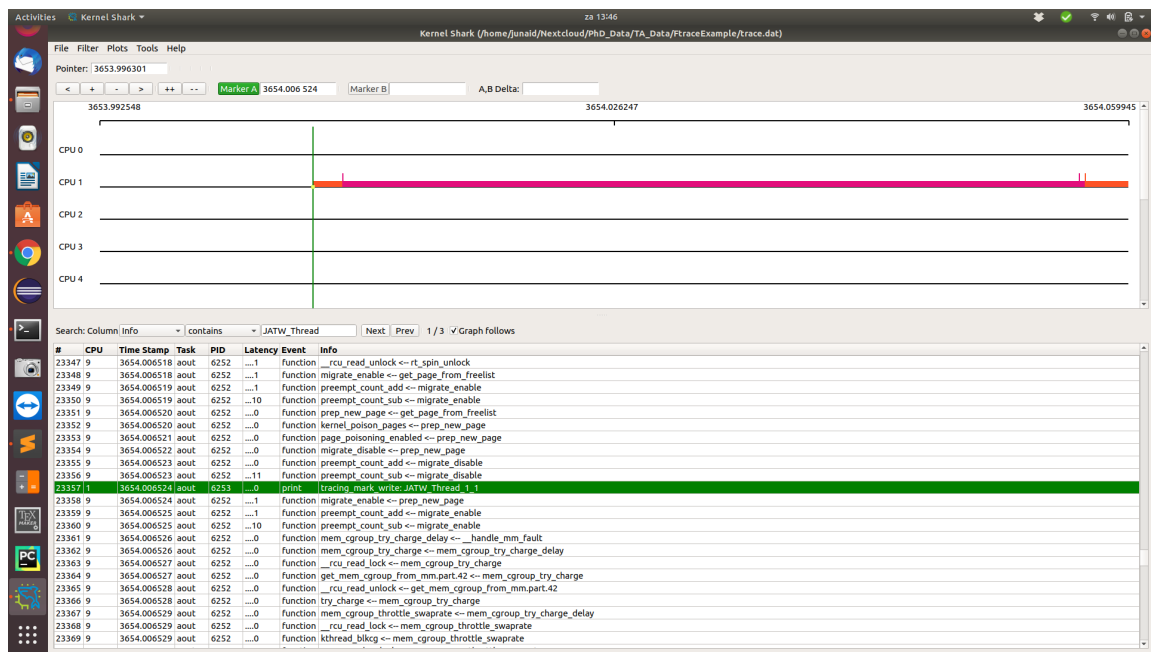


Figure 10: Trace text search box